

```
# import libraries
```

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set()
import warnings
warnings.filterwarnings("ignore")
```

```
dataset = pd.read_csv("CarPrice_Assignment.csv")
dataset.head()
```

	car_ID	symboling	CarName	fueltype	aspiration
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std

	carbody	drivewheel	engine	location	wheelbase	...
0	convertible	rwd	front	88.6	...	130
1	convertible	rwd	front	88.6	...	130
2	hatchback	rwd	front	94.5	...	152
3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136

	fuelsystem	bore	ratio	stroke	compressionratio	horsepower	peakrpm
0	mpfi	3.47	2.68	9.0	111	5000	
1	mpfi	3.47	2.68	9.0	111	5000	
2	mpfi	2.68	3.47	9.0	154	5000	
3	mpfi	3.19	3.40	10.0	102	5500	

```
24
4      mpfi      3.19      3.40      8.0      115      5500
18
```

```
      highwaympg      price
0           27  13495.0
1           27  16500.0
2           26  16500.0
3           30  13950.0
4           22  17450.0
```

```
[5 rows x 26 columns]
```

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 26 columns):
```

#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

```
dtypes: float64(8), int64(8), object(10)
```

```
memory usage: 41.8+ KB
```

```
dataset.duplicated().sum()
```

0

```
dataset = dataset.drop(['CarName'], axis=1)
```

```
dataset.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth
carheight \					
count	205.000000	205.000000	205.000000	205.000000	205.000000
205.000000					
mean	103.000000	0.834146	98.756585	174.049268	65.907805
53.724878					
std	59.322565	1.245307	6.021776	12.337289	2.145204
2.443522					
min	1.000000	-2.000000	86.600000	141.100000	60.300000
47.800000					
25%	52.000000	0.000000	94.500000	166.300000	64.100000
52.000000					
50%	103.000000	1.000000	97.000000	173.200000	65.500000
54.100000					
75%	154.000000	2.000000	102.400000	183.100000	66.900000
55.500000					
max	205.000000	3.000000	120.900000	208.100000	72.300000
59.800000					

	curbweight	enginesize	boreratio	stroke
compressionratio \				
count	205.000000	205.000000	205.000000	205.000000
205.000000				
mean	2555.565854	126.907317	3.329756	3.255415
10.142537				
std	520.680204	41.642693	0.270844	0.313597
3.972040				
min	1488.000000	61.000000	2.540000	2.070000
7.000000				
25%	2145.000000	97.000000	3.150000	3.110000
8.600000				
50%	2414.000000	120.000000	3.310000	3.290000
9.000000				
75%	2935.000000	141.000000	3.580000	3.410000
9.400000				
max	4066.000000	326.000000	3.940000	4.170000
23.000000				

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000

50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

Checking missing values

```
dataset.isnull().sum()
```

```
car_ID          0
symboling       0
fueltype        0
aspiration      0
doornumber      0
carbody         0
drivewheel      0
enginelocation  0
wheelbase       0
carlength       0
carwidth        0
carheight       0
curbweight      0
enginetype      0
cylindernumber  0
enginesize      0
fuelsystem      0
boreratio       0
stroke          0
compressionratio 0
horsepower      0
peakrpm        0
citympg         0
highwaympg      0
price           0
dtype: int64
```

we have to handle outlier for the listed variables
wheelbase, carlength, carwidth, enginesize, stroke,
compressionratio, horsepower, peakrpm, citympg, highwaympg

```
columns_to_check = ['wheelbase', 'carlength', 'carwidth',
                    'enginesize', 'stroke',
                    'compressionratio', 'horsepower', 'peakrpm',
                    'citympg', 'highwaympg']
```

Function to detect and treat outliers

```
def treat_outliers(dataset, columns):
    for col in columns:
        # Calculate Q1, Q3, and IQR
        Q1 = dataset[col].quantile(0.25)
        Q3 = dataset[col].quantile(0.75)
```

```

IQR = Q3 - Q1

# Define lower and upper bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Cap the outliers
dataset[col] = dataset[col].apply(lambda x: upper_bound if x >
upper_bound else (lower_bound if x < lower_bound else x))

return dataset


# Apply the function
dataset = treat_outliers(dataset, columns_to_check)


def boxplots(col):
    sns.boxplot(dataset[col])
    plt.show()


for i in list(dataset.select_dtypes(exclude=['object']).columns)[0:]:
    boxplots(i)



```
























```
#Convert categorical features (fueltype, aspiration, doornumber,
carbody, drivewheel, enginelocation,
#enginetype, cylindernumber, fuelsystem)
#into numerical using Label Encoding.
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
dataset['fueltype']= le.fit_transform(dataset['fueltype'])
dataset['aspiration']= le.fit_transform(dataset['aspiration'])
dataset['doornumber']= le.fit_transform(dataset['doornumber'])
dataset['carbody']= le.fit_transform(dataset['carbody'])
dataset['drivewheel']= le.fit_transform(dataset['drivewheel'])
dataset['enginelocation']= le.fit_transform(dataset['enginelocation'])
dataset['enginetype']= le.fit_transform(dataset['enginetype'])
dataset['cylindernumber']= le.fit_transform(dataset['cylindernumber'])
dataset['fuelsystem']= le.fit_transform(dataset['fuelsystem'])
```

```
dataset.head()
```

	car_ID	symboling	fueltype	aspiration	doornumber	carbody
drivewheel						
0	1	3	1	0	1	0
2						
1	2	3	1	0	1	0
2						
2	3	1	1	0	1	2
2						

3	4	2	1	0	0	3
1						
4	5	2	1	0	0	3
0						

	engine	location	wheelbase	carlength	...	enginesize
fuelsystem \						
0		0	88.6	168.8	...	130.0
1		0	88.6	168.8	...	130.0
2		0	94.5	171.2	...	152.0
3		0	99.8	176.6	...	109.0
4		0	99.4	176.6	...	136.0

	boreratio	stroke	compressionratio	horsepower	peakrpm
citympg \					
0	3.47	2.68	9.0	111.0	5000.0
1	3.47	2.68	9.0	111.0	5000.0
2	2.68	3.47	9.0	154.0	5000.0
3	3.19	3.40	10.0	102.0	5500.0
4	3.19	3.40	8.0	115.0	5500.0

	highwaympg	price
0	27.0	13495.0
1	27.0	16500.0
2	26.0	16500.0
3	30.0	13950.0
4	22.0	17450.0

[5 rows x 25 columns]

```
x = dataset.iloc[:,0:-1] # Independent variable
y = dataset.iloc[:, -1] # Dependent variable (Price)
```

```
x.head()
```

	car_ID	symboling	fueltype	aspiration	doornumber	carbody
drivewheel \						
0	1	3	1	0	1	0
2						
1	2	3	1	0	1	0
2						

2	3	1	1	0	1	2
2						
3	4	2	1	0	0	3
1						
4	5	2	1	0	0	3
0						

engine	location	wheelbase	carlength	...	cylinders	number
0	130.0	0	88.6	168.8	...	2
1	130.0	0	88.6	168.8	...	2
2	152.0	0	94.5	171.2	...	3
3	109.0	0	99.8	176.6	...	2
4	136.0	0	99.4	176.6	...	1

fuel	system	bore	ratio	stroke	compression	ratio	horsepower
0	5000.0	5	3.47	2.68	9.0	111.0	
1	5000.0	5	3.47	2.68	9.0	111.0	
2	5000.0	5	2.68	3.47	9.0	154.0	
3	5500.0	5	3.19	3.40	10.0	102.0	
4	5500.0	5	3.19	3.40	8.0	115.0	

city	mpg	highway	mpg
0	21.0	27.0	
1	21.0	27.0	
2	19.0	26.0	
3	24.0	30.0	
4	18.0	22.0	

[5 rows x 24 columns]

y.head()

0	13495.0
1	16500.0
2	16500.0
3	13950.0
4	17450.0

Name: price, dtype: float64

Feature Scaling

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
sc_x = scaler.fit_transform(x)
pd.DataFrame(sc_x)
```

	0	1	2	3	4	5	
6 \							
0	-1.723622	1.743470	0.328798	-0.469295	1.130388	-3.050975	
1.213330							
1	-1.706724	1.743470	0.328798	-0.469295	1.130388	-3.050975	
1.213330							
2	-1.689826	0.133509	0.328798	-0.469295	1.130388	-0.717207	
1.213330							
3	-1.672928	0.938490	0.328798	-0.469295	-0.884652	0.449677	
0.589081							
4	-1.656029	0.938490	0.328798	-0.469295	-0.884652	0.449677	
2.391492							
..	
...							
200	1.656029	-1.476452	0.328798	-0.469295	-0.884652	0.449677	
1.213330							
201	1.672928	-1.476452	0.328798	2.130854	-0.884652	0.449677	
1.213330							
202	1.689826	-1.476452	0.328798	-0.469295	-0.884652	0.449677	
1.213330							
203	1.706724	-1.476452	-3.041381	2.130854	-0.884652	0.449677	
1.213330							
204	1.723622	-1.476452	0.328798	2.130854	-0.884652	0.449677	
1.213330							
	7	8	9	...	14	15	
16 \							
0	-0.121867	-1.723005	-0.426521	...	-0.147475	0.160196	0.869568
1	-0.121867	-1.723005	-0.426521	...	-0.147475	0.160196	0.869568
2	-0.121867	-0.717590	-0.231513	...	1.112210	0.809329	0.869568
3	-0.121867	0.185580	0.207256	...	-0.147475	-0.459430	0.869568
4	-0.121867	0.117416	0.207256	...	-1.407161	0.337232	0.869568
..
200	-0.121867	1.770387	1.198549	...	-0.147475	0.484762	0.869568
201	-0.121867	1.770387	1.198549	...	-0.147475	0.484762	0.869568
202	-0.121867	1.770387	1.198549	...	1.112210	1.428955	0.869568


```

203 -0.121867  1.770387  1.198549  ...  1.112210  0.602787 -0.126306
204 -0.121867  1.770387  1.198549  ... -0.147475  0.484762  0.869568

```

```

          17          18          19          20          21          22
23
0    0.519071 -2.106623 -0.049433  0.229801 -0.262757 -0.649321 -
0.552143
1    0.519071 -2.106623 -0.049433  0.229801 -0.262757 -0.649321 -
0.552143
2   -2.404880  0.753841 -0.049433  1.441341 -0.262757 -0.958163 -
0.702161
3   -0.517266  0.500383  1.214121 -0.023777  0.791357 -0.186058 -
0.102086
4   -0.517266  0.500383 -1.312986  0.342502  0.791357 -1.112584 -
1.302237
..          ...          ...          ...          ...          ...
...
200  1.666445 -0.404828  0.582344  0.314327  0.580534 -0.340479 -
0.402124
201  1.666445 -0.404828 -0.428499  1.610393  0.369711 -0.958163 -
0.852180
202  0.926204 -1.418663 -0.302143  0.877834  0.791357 -1.112584 -
1.152218
203 -1.183483  0.500383  1.972253  0.088924 -0.684403  0.122783 -
0.552143
204  1.666445 -0.404828  0.582344  0.314327  0.580534 -0.958163 -
0.852180

```

[205 rows x 24 columns]

```
pd.DataFrame(sc_x).describe()
```

```

          0          1          2          3
4  \
count  2.050000e+02  2.050000e+02  2.050000e+02  2.050000e+02
2.050000e+02
mean   -6.932124e-17  4.332578e-17 -7.798640e-17  6.282238e-17
1.039819e-16
std     1.002448e+00  1.002448e+00  1.002448e+00  1.002448e+00
1.002448e+00
min    -1.723622e+00 -2.281433e+00 -3.041381e+00 -4.692953e-01 -
8.846517e-01
25%    -8.618111e-01 -6.714717e-01  3.287980e-01 -4.692953e-01 -
8.846517e-01
50%     0.000000e+00  1.335090e-01  3.287980e-01 -4.692953e-01 -
8.846517e-01
75%     8.618111e-01  9.384897e-01  3.287980e-01 -4.692953e-01

```

```

1.130388e+00
max    1.723622e+00  1.743470e+00  3.287980e-01  2.130854e+00
1.130388e+00

```

	5	6	7	8
9 \				
count	2.050000e+02	2.050000e+02	2.050000e+02	2.050000e+02
mean	-1.646380e-16	-1.126470e-16	-8.665155e-18	-1.065814e-15
std	1.002448e+00	1.002448e+00	1.002448e+00	1.002448e+00
min	-3.050975e+00	-2.391492e+00	-1.218667e-01	-2.063824e+00
25%	-7.172069e-01	-5.890807e-01	-1.218667e-01	-7.175899e-01
50%	4.496773e-01	-5.890807e-01	-1.218667e-01	-2.915663e-01
75%	4.496773e-01	1.213330e+00	-1.218667e-01	6.286445e-01
max	1.616562e+00	1.213330e+00	8.205689e+00	2.647996e+00

	...	14	15	16	17 \
count	...	2.050000e+02	2.050000e+02	2.050000e+02	2.050000e+02
mean	...	-1.386425e-16	-1.992986e-16	1.039819e-16	2.252940e-15
std	...	1.002448e+00	1.002448e+00	1.002448e+00	1.002448e+00
min	...	-2.666846e+00	-1.875720e+00	-1.620116e+00	-2.923049e+00
25%	...	-1.474754e-01	-8.135028e-01	-1.122179e+00	-6.653141e-01
50%	...	-1.474754e-01	-1.348641e-01	8.695675e-01	-7.312136e-02
75%	...	-1.474754e-01	4.847625e-01	8.695675e-01	9.262039e-01
max	...	4.891266e+00	2.432160e+00	1.865441e+00	2.258638e+00

	18	19	20	21
22 \				
count	2.050000e+02	2.050000e+02	2.050000e+02	2.050000e+02
mean	2.980813e-15	6.845473e-16	-2.057974e-17	-6.932124e-16
std	1.002448e+00	1.002448e+00	1.002448e+00	1.002448e+00
min	-2.179040e+00	-2.071118e+00	-1.545246e+00	-2.054752e+00
25%	-5.496614e-01	-5.548541e-01	-9.253886e-01	-6.844030e-01
50%	1.020901e-01	-4.943268e-02	-2.210047e-01	1.588884e-01
75%	5.365910e-01	4.559888e-01	3.706777e-01	7.913570e-01
max	2.165970e+00	1.972253e+00	2.314777e+00	3.004997e+00

3.288412e+00

```
count    23
count    2.050000e+02
mean     9.856614e-17
std      1.002448e+00
min     -2.202350e+00
25%     -8.521802e-01
50%     -1.020860e-01
75%      4.979894e-01
max      2.523244e+00
```

[8 rows x 24 columns]

```
# It's multiple linear regression, hence we have to check
"Multicollinearity"
variable = sc_x
variable.shape

(205, 24)
```

To check Multicollinearity - VIF (Variance Inflation Factor)

```
# VIF value is more than 5 means there is a multicollinearity there in
the features
from statsmodels.stats.outliers_influence import
variance_inflation_factor
variable = sc_x
vif = pd.DataFrame()
vif['variance_inflation_factor'] =
[variance_inflation_factor(variable, i) for i in
range(variable.shape[1])]


vif['Features'] = x.columns
vif
```

	variance_inflation_factor	Features
0	1.479131	car_ID
1	2.828315	symboling
2	5.584468	fueltype
3	3.122923	aspiration
4	2.780304	doornumber
5	2.813480	carbody
6	2.626872	drivewheel
7	1.805014	engine location
8	10.551627	wheelbase
9	11.537443	carlength
10	8.040508	carwidth
11	2.929516	carheight

12	18.449284	curbweight
13	1.442345	enginetype
14	2.204296	cylindernumber
15	16.377475	enginesize
16	2.757890	fuelsystem
17	2.491841	boreratio
18	1.697226	stroke
19	4.315996	compressionratio
20	17.445684	horsepower
21	2.970193	peakrpm
22	23.176338	citympg
23	22.110875	highwaympg

Finding correlation

```
plt.figure(figsize=(20,15))
#corr = Housing.corr()
sns.heatmap(dataset.corr(), annot=True, cmap='coolwarm')
plt.show()
```



We have multicorrelated variables hence we need to apply PCA for dimensionality reduction

```
# Apply PCA
from sklearn.decomposition import PCA

pca = PCA(n_components=0.95) # Retain 95% of the variance
X_pca = pca.fit_transform(sc_x)

# Let's check the explained variance ratio to understand how much
information each principal component retains

explained_variance = pca.explained_variance_ratio_
explained_variance

array([0.32886298, 0.16703802, 0.08388498, 0.06767003, 0.05224637,
       0.04496342, 0.04141339, 0.03759184, 0.02859023, 0.02565239,
       0.02152603, 0.01876589, 0.01566118, 0.01416775, 0.01235966])
```

```
# Split the Data into Train and Test Sets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_pca, y,
test_size=0.2, random_state=42)
```

Linear Regression

```
# Apply Linear Regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Initialize the Linear Regression model
lr_model = LinearRegression()

# Train the model
lr_model.fit(X_train, y_train)

# Make predictions
y_pred_lr = lr_model.predict(X_test)
y_pred_lr

array([24495.64959172, 21391.61915098,  9280.41659459, 14280.61097737,
      26711.1821066 ,  6593.44981859,  7609.43157162,  6705.47725175,
      10787.37603735,  6103.55009794, 15386.77862939,  7481.39209327,
      15623.12078086, 11115.45256861, 33341.20982346,  4873.97687978,
     -4220.80495503, 14118.81572874,  9798.09584758, 11299.01521206,
      11299.17916788, 21308.15377546,  6883.26415079, 1776.51983521,
      6805.8324699 , 25958.839732 , 13889.49843904, 15831.3596836 ,
      4929.01094257, 17103.50952177, 27360.17456389,  6120.31202415,
      4707.26111497, 22375.70714452,  7819.91713441, 28456.93563104,
      10164.205043 , 10447.43562098,  5516.22217562, 14293.09217264,
      9136.07052383])

# Evaluate the Linear Regression model
mae_lr = mean_absolute_error(y_test, y_pred_lr)
mse_lr = mean_squared_error(y_test, y_pred_lr)
r2_lr = r2_score(y_test, y_pred_lr)

print("Linear Regression Evaluation:")
print(f"MAE: {mae_lr}")
print(f"MSE: {mse_lr}")
print(f"R-squared: {r2_lr}")

Linear Regression Evaluation:
MAE: 2664.127335130712
MSE: 15831151.943431232
R-squared: 0.7994635721992898
```

Random Forest Regression

```
# Apply Random Forest Regression
from sklearn.ensemble import RandomForestRegressor

# Initialize the Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Train the model
rf_model.fit(X_train, y_train)

# Make predictions
y_pred_rf = rf_model.predict(X_test)
y_pred_rf

array([21234.63 , 18537.87 ,  8714.31 , 13026.23 , 25965.325,
        6058.375,
        7265.    ,  7981.51 , 11639.49 ,  7341.47 , 16354.65 ,  7795.71
,
        20358.115, 11458.66 , 29814.24 ,  6725.2  ,  6060.105, 14912.97
,
        8046.03 , 11253.13 , 10940.97 , 14698.61 ,  5919.275,  6733.36
,
        7302.65 , 26782.405,  9648.25 , 16688.52 ,  7312.43 , 16223.62
,
        30686.165,  7150.98 ,  7857.12 , 19450.63 ,  8308.2  ,
        27680.455,
        11402.46 , 11931.18 ,  6975.285, 15020.66 ,  8733.14 ]))


# Evaluate the Random Forest model
mae_rf = mean_absolute_error(y_test, y_pred_rf)
mse_rf = mean_squared_error(y_test, y_pred_rf)
r2_rf = r2_score(y_test, y_pred_rf)

print("\nRandom Forest Regression Evaluation:")
print(f"MAE: {mae_rf}")
print(f"MSE: {mse_rf}")
print(f"R-squared: {r2_rf}")


Random Forest Regression Evaluation:
MAE: 2090.739463414634
MSE: 13142203.156256925
R-squared: 0.8335250344507925

# Prediction vs True Values Plot (for Random Forest)
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_rf, color='blue', alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         color='red', linestyle='--')
plt.title('Random Forest: Prediction vs True Values')
```

```
plt.xlabel('True Values')
plt.ylabel('Predicted Values')
plt.show()
```



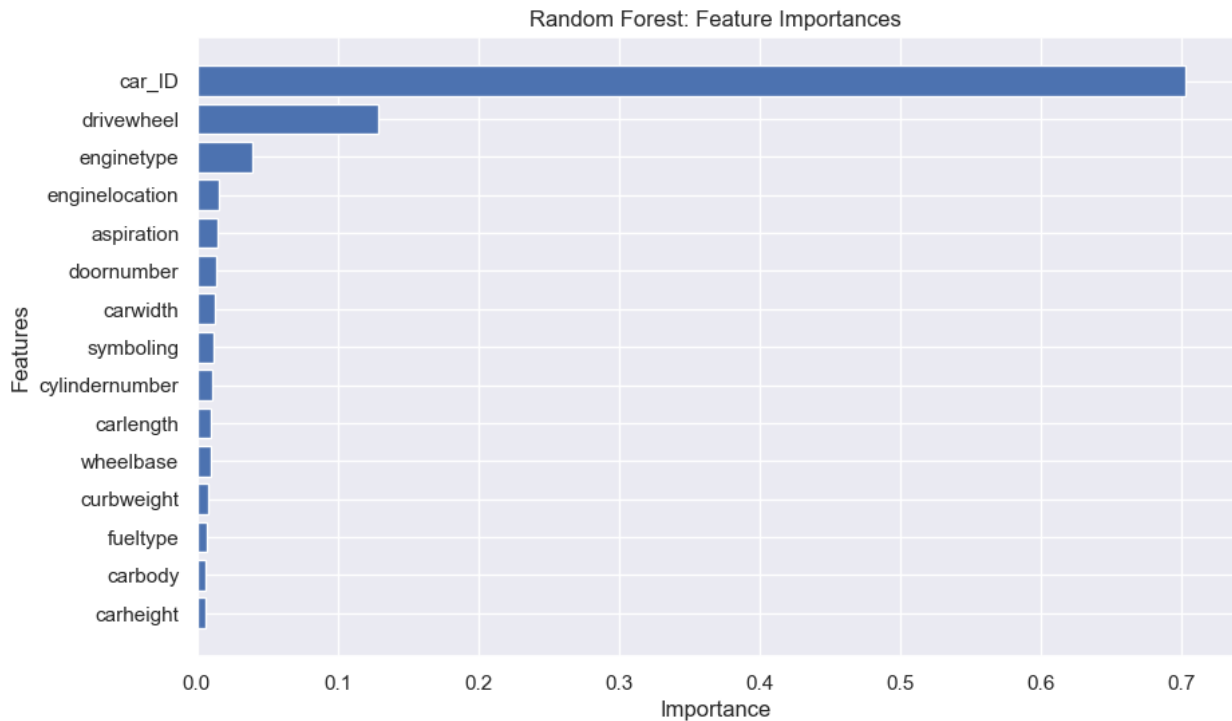
```
# Residuals Plot (Random Forest)
residuals_rf = y_test - y_pred_rf
plt.figure(figsize=(10, 6))
sns.residplot(x=y_pred_rf, y=residuals_rf, lowess=True, color='blue',
line_kws={'color': 'red', 'lw': 1})
plt.title('Random Forest: Residuals Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()
```

```
# Feature Importance (Random Forest)
importances = rf_model.feature_importances_
features = x.columns

# Sort feature importances in descending order
indices = importances.argsort()

plt.figure(figsize=(10, 6))
plt.barh(range(len(indices)), importances[indices], align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.title('Random Forest: Feature Importances')
plt.xlabel('Importance')
plt.ylabel('Features')
plt.show()
```



Ridge Regression

```
from sklearn.linear_model import Ridge, Lasso

# Ridge Regression
ridge_model = Ridge(alpha=1.0)
ridge_model.fit(X_train, y_train)

# Make predictions and evaluate Ridge
y_pred_ridge = ridge_model.predict(X_test)
print("Ridge Regression Performance:")
print(f"MAE: {mean_absolute_error(y_test, y_pred_ridge)}")
print(f"MSE: {mean_squared_error(y_test, y_pred_ridge)}")
print(f"R-squared: {r2_score(y_test, y_pred_ridge)}")
```

```
Ridge Regression Performance:
MAE: 2660.1791265483457
MSE: 15824634.958824884
R-squared: 0.7995461241713546
```

Lasso Regression

```
# Lasso Regression
lasso_model = Lasso(alpha=0.1)
lasso_model.fit(X_train, y_train)

# Make predictions and evaluate Lasso
y_pred_lasso = lasso_model.predict(X_test)
```

```
print("\nLasso Regression Performance:")
print(f"MAE: {mean_absolute_error(y_test, y_pred_lasso)}")
print(f"MSE: {mean_squared_error(y_test, y_pred_lasso)}")
print(f"R-squared: {r2_score(y_test, y_pred_lasso)}")
```

Lasso Regression Performance:
MAE: 2664.0331372881587
MSE: 15830174.791726185
R-squared: 0.7994759499790653

Gradient Boosting Regressor

```
from sklearn.ensemble import GradientBoostingRegressor

# Apply Gradient Boosting Regressor
gb_model = GradientBoostingRegressor(n_estimators=100,
learning_rate=0.1, max_depth=3, random_state=42)
gb_model.fit(X_train, y_train)

# Make predictions and evaluate GB
y_pred_gb = gb_model.predict(X_test)
print("\nGradient Boosting Regression Performance:")
print(f"MAE: {mean_absolute_error(y_test, y_pred_gb)}")
print(f"MSE: {mean_squared_error(y_test, y_pred_gb)}")
print(f"R-squared: {r2_score(y_test, y_pred_gb)}")
```

Gradient Boosting Regression Performance:
MAE: 2151.261016339273
MSE: 13214894.044313649
R-squared: 0.8326042441585482

XGBoost Regressor

```
import xgboost as xgb

# Apply XGBoost Regressor
xgb_model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1,
max_depth=3, random_state=42)
xgb_model.fit(X_train, y_train)

# Make predictions and evaluate XGBoost
y_pred_xgb = xgb_model.predict(X_test)
print("\nXGBoost Regression Performance:")
print(f"MAE: {mean_absolute_error(y_test, y_pred_xgb)}")
print(f"MSE: {mean_squared_error(y_test, y_pred_xgb)}")
print(f"R-squared: {r2_score(y_test, y_pred_xgb)}")
```

XGBoost Regression Performance:

MAE: 1991.606135480183
MSE: 11071641.628366638
R-squared: 0.8597532592716014